

An aerial, black and white photograph of the Amazon River delta. The river branches out into numerous smaller channels and floodplains. A large, elongated island covered in dense forest is prominent in the foreground, with the river flowing alongside it. The sky is filled with large, dramatic clouds, and a bright light source, possibly the sun, is visible on the horizon, creating a strong contrast and long shadows across the landscape.

Understanding the Amazon from Space

Louis Limnavong, Diederik Ketellapper, Daniel Wientjens, Alessandro Girelli, Eric Brea Garcia and Pierre Watrin

AGENDA

Project Overview

Classification of weather conditions

Object classification



Understanding the object classification: Class activation maps

Q&A

THE CHALLENGE

To label satellite image chips with atmospheric conditions and various classes of land cover/land use. Resulting algorithms will help the global community better understand where, how, and why deforestation happens all over the world - and ultimately how to respond.

Working with multi-labeled satellite images



AGENDA

Project Overview



Classification of weather conditions

Object classification

Understanding the object classification: Class activation maps


Q&A

First we focused on the weather classification



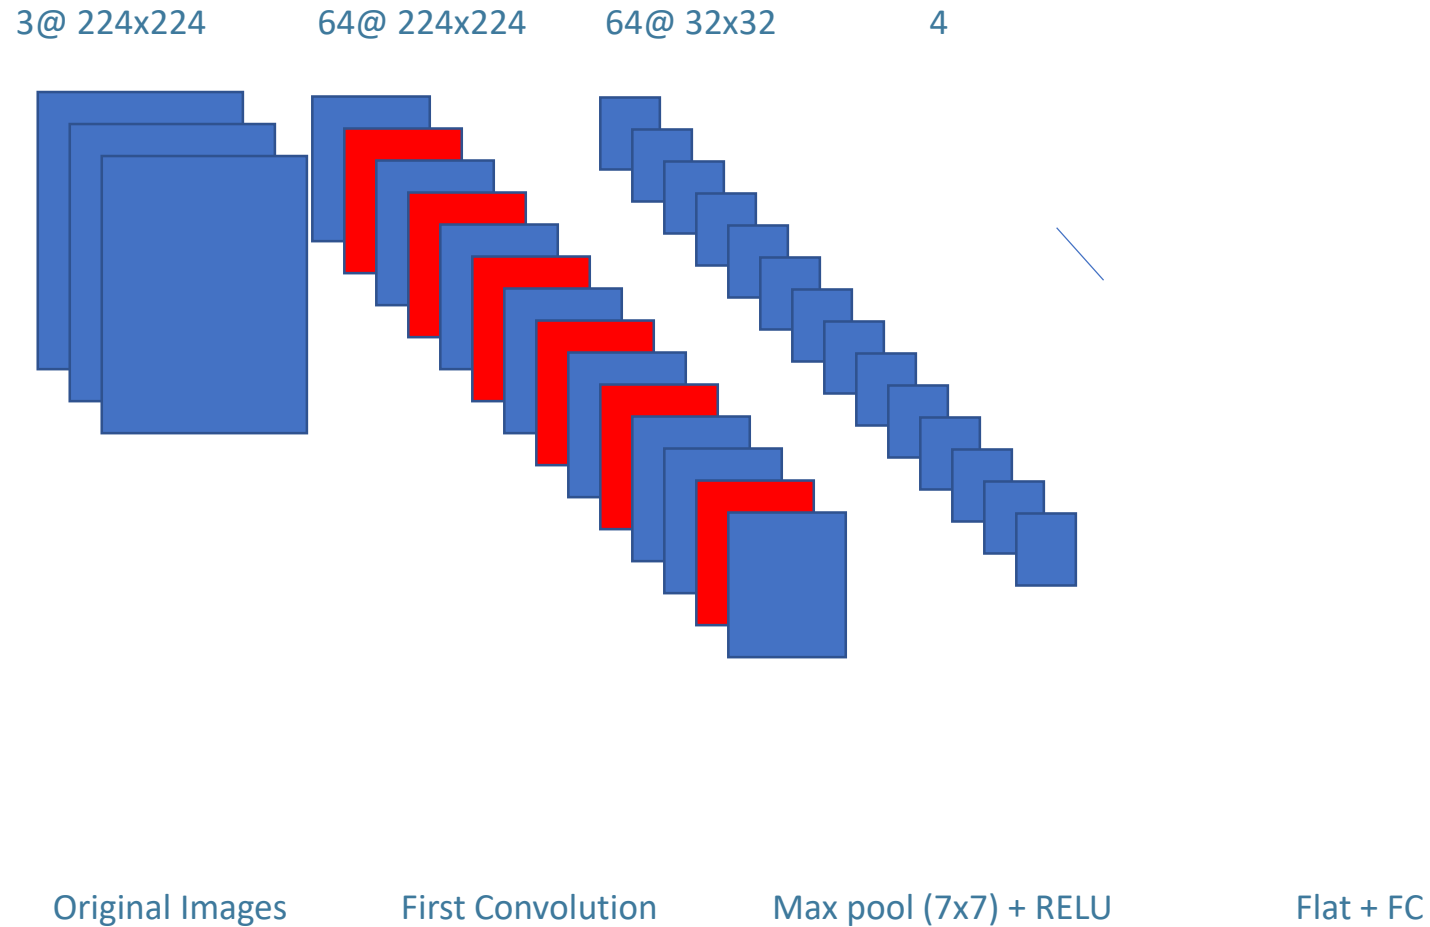
There are four weather related labels, which are mutually exclusive

Starting from scratch, a single-layer CNN




- The model was trained on 6400 images for 20 epochs
- The training accuracy was 0.87
- The testing accuracy was 0.67
- The model is performing worse in the test set, than in the train set, which means is **overfitting** the data

Then we add a dropout layer to reduce overfitting



- The model was trained on 6400 images for 20 epochs
- The training accuracy was 0.73
- The testing accuracy was 0.68
- We are able to reduce overfitting, but that doesn't translate into an improved test accuracy


Finally, we use a deeper architecture




We achieved an accuracy of 0.80 in train, and 0.78 in test.

Comparison across models

Loss



Accuracy



AGENDA

Project Overview



Classification of weather conditions

Object classification

Understanding the object classification: Class activation maps

Q&A

Now we identify roads and water bodies



Among the clear images we classify the presence of roads and water

We fine tuned Resnet18 for binary classification

```
model_resnet = models.resnet18(pretrained=True)

for param in model_resnet.parameters():
    param.requires_grad = False
model_resnet.fc = nn.Linear(512, 2, bias = True)
```

- The model was trained on 3200 images for 15 epochs, and tested it on 800
- It achieved a 0.8409 accuracy on the train set and 0.8425 on the test set

AGENDA

Project Overview


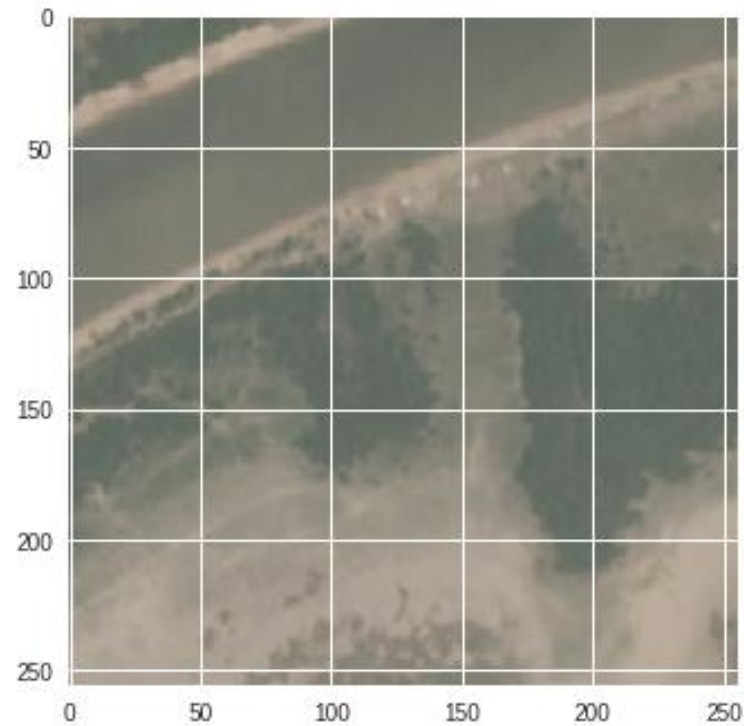
Classification of weather conditions

Object classification

Understanding the object classification: Class activation maps



Q&A

Plotting the discriminative image regions used by the CNN



The model uses the activated areas (red) to predict the probability of having water in the image

Plotting the discriminative image regions used by the CNN



The model uses the activated areas (red) to predict the probability of having road in the image

AGENDA

Project Overview

Classification of weather conditions

Object classification

Understanding the object classification: Class activation maps

Q&A
